

Learning Sub-Grid Physics of Geophysical Turbulence with Fourier Neural Operators: A Case Study in Quasi-Biennial Oscillation Modelling

YanJun XUE

May 27, 2026

Abstract

Eddy parameterization is known to play a fundamental role in earth system models. The Fourier Neural Operator (FNO) is a newly-proposed deep learning algorithm specialized in learning PDEs. Here we use the gravity wave parameterization in the 1D-QBO model as a testbed for the application of FNOs in eddy parameterization. An FNO with a very small parameter size can simulate the QBO period more accurately than a deep MLP in this task. However, FNOs still have difficulty learning the wave forcing with a wide wave spectrum. These findings demonstrate the potential of FNOs over traditional MLPs in gravity wave parameterization of QBO and provide a possible way for integrating AI-based schemes into complex global circulation models.

1 Introduction

Ocean and climate models simulate the dynamics of the geophysical flow by solving the associated differential equations numerically. For computational representation, the continuous fields of physical variables (e.g., velocity, pressure, temperature) in the atmosphere and ocean must be discretized. The most intuitive approach is to divide the fluid into a number of grid cells and approximate the continuous field by the average value across the grid cell or the value at the center of the grid cell (Neelin [2010]). This can approximately capture the behavior of motions at spacial scales much larger than the grid cell, but obviously omits the other values at different points within the grid cell. Therefore, higher resolution models enable more detailed simulations and explicitly captures the complexities of physical systems. However, if we try to increase the spatial resolution, the consumption of computation cost will grow exponentially. Taking 2D ocean modeling as an example, a doubling of the horizontal resolution results in an approximately four-fold increase in the number of grid points. Additionally, according to the CFL condition, the time step should be shortened correspondingly to maintain numerical stability (De Moura and Kubrusly [2013]). Thus, there is a tradeoff between physical accuracy and computing cost.

To bridge the aforementioned gap, one strategy is to employ a coarse-resolution model and construct functions of resolved variables to parameterize the processes whose spatial scales are smaller than the grid cell, which are called subgrid-scale parameterizations (Uliasz [1994]). Since geophysical turbulence is inherently multi-scale (McWilliams [1989]), among different types of subgrid parameterization, eddy parameterization has consistently remained a critical topic in both ocean and climate modeling (Fox-Kemper et al. [2019]; Bauer et al. [2015]). In the oceans, mesoscale eddies at 10^1 to 10^2 km play a crucial role in the global transport and distribution of heat, salt, nutrients, and carbon dioxide, while they cannot be fully resolved in the most earth system models at 1° or 0.25° resolution (Hewitt et al. [2020]). Actually, the model effect of eddy representation is primarily governed by the ratio of the baroclinic deformation radius to the resolution (Hallberg [2013]). As a result, the resolution in the climate models can be a bit coarser due to the larger deformation radius in the atmosphere. However, the climate model grid size is about 2° in the horizontal, still not enough for representing some subgrid effects because the atmospheric models are more expensive for some extra parameterization schemes, such as radiation, convection and clouds (Neelin [2010]). Therefore, the eddy parameterization is both

necessary in ocean and climate models.

Gravity Wave Parameterization (GWP) is one important component in the atmospheric models. Gravity waves in the atmosphere are characterized by their high frequency and short wavelength, so some of those with even shorter wavelengths cannot be represented in the numerical models. Gravity waves are the momentum sources for the large-scale motions in the atmosphere, thus affecting the long-term climate predictions significantly. Quasi-Biennial Oscillation (QBO) is one of the atmospheric signals that are driven by gravity waves (Baldwin et al. [2001]). The mechanism of QBO can be explained by wave-mean flow interaction theory (Plumb [1977]). The details of GWP schemes are discussed in section 2. Additionally, realistic GWP is crucial for simulating the QBO in Global Circulation Models (GCMs, Richter et al. [2020]), and consequently affects other signals in the atmosphere (such as MJO and subtropical jets) through teleconnections, which are closely related to subseasonal-to-seasonal forecast skills (Anstey et al. [2022a]). Garfinkel et al. [2022]’s study has shown that while the QBO’s period and amplitude can be effectively tuned through parameterized gravity waves, its vertical structure remains poorly captured and difficult to simulate accurately in the GCMs. Thus, we adopt the QBO as a testbed to study the general geophysical turbulence parameterization.

Over the past decade, Machine Learning (ML) methods have been increasingly applied to the eddy parameterization in the atmospheric and oceanic sciences (de Burgh-Day and Leeuwenburg [2023]). Deep neural networks have demonstrated their potential to parameterize convection processes in climate models (Rasp et al. [2018]) and 2D turbulence in ocean models (Yan et al. [2025]). The current deficiencies of ML-driven parameterization primarily lie in the physical inconsistency and their ability to generalize. ML models cannot learn the conservation laws directly from training data alone and thus need extra physical constraints (Beucler et al. [2021]). Regarding generalization, O’Gorman and Dwyer [2018]’s study also found that global warming can be accurately simulated only when training samples from both the control and warm climate are used. Existing ML strategies for GWP in QBO mainly focus on convolutional neural networks (CNNs, Pahlavan et al. [2024]), whose performances are sensitive to the receptive field (Pahlavan et al. [2025]). Recently, Fourier Neural Operator (FNO) was proposed as an operator learning method and has shown strong capability in modeling the Navier–Stokes equations and other differential equations-based problems (Li et al. [2020]). The results suggest FNOs have potential for sub-grid turbulence parameterization, which is also characterized by intense nonlinearity and complex interactions. Therefore, in this study, we will investigate the capability of FNO for GWP within a 1D-QBO testbed. In Section 2, we introduce the 1D-QBO model we employ and the model architecture of FNO in detail. Section 3 describes our results, and the discussion and summary are provided in Section 4.

2 Methods

2.1 1D-QBO Model

QBO is the dominant large-scale interannual variability in the tropical stratosphere. It is characterized by downward propagating reversals of the mean zonal winds, with a period of approximately 28 months (Baldwin et al. [2001]). As shown in Figure 1(a), the mean zonal winds over Singapore obtained from radiosonde measurement demonstrate alternating easterly and westerly regimes that descend through the stratosphere as a function of time. The FFT spectrum of the zonal wind at 35.5 km shown in Figure 1(b) further reveals a sharp peak near the quasi-biennial timescale of 28 months.

The persuasive explanation of QBO was first proposed by Lindzen and Holton [1968]. They regarded QBO as the absorption of gravity waves at critical layers, which was inspired by the general critical layer theory (Booker and Bretherton [1967]). In this perspective, the momentum sources of the zonal mean winds acceleration comes from the upward propagating gravity waves excited at the top of the troposphere, including Kelvin wave and Rossby-gravity wave. If the background zonal wind approaches the phase speed of these gravity waves, the gravity waves will be refracted, leading to a strong convergence of wave momentum flux at this level, which changes the mean flows. This




Figure 1: (a) Monthly zonal mean wind over Singapore and (b) its FFT spectrum at 35.5 km.

mechanism can be described by the following mean flow equation,

$$\rho_0 \left(\frac{\partial \bar{u}}{\partial t} + w \frac{\partial \bar{u}}{\partial z} \right) = - \frac{dF_{MW}}{dz} = -f(\bar{u}) \frac{\partial \bar{u}}{\partial z}, \quad (1)$$

where Lindzen parameterized the gravity wave forcing (GWF) as a nonlinear function of zonal wind shear. If the convection term is neglected, Equation 1 explains the downward propagation of QBO phase with a speed of $f(\bar{u})/\rho_0$. In summary, the mechanism of Lindzen-Holton model involves two processes. The background flows first modulate the way waves propagate and thus lead to the convergence of wave momentum flux. Then, the wave exerts a force on and alters the mean flow.

Holton and Lindzen [1972] and Plumb [1977] further developed this theory and showed that just two opposite wave modes are required to produce a QBO phase if mean flow diffusion exists. This approach provided a more rigorous and practical model to study QBO, which was called 1D-QBO model,

$$\frac{\partial \bar{u}}{\partial t} + w \frac{\partial \bar{u}}{\partial z} - \kappa \frac{\partial^2 \bar{u}}{\partial z^2} = \text{GWF}(\bar{u}, z), \quad (2)$$

where,

$$\text{GWF}(\bar{u}, z) = - \frac{\rho_L}{\rho_0} \sum_n \frac{\partial F_n}{\partial z}; \quad F_n = F_L \exp \left\{ - \int_{z_L}^z \frac{N\mu}{k_n [\bar{u}(z) - c_n]^2} dz' \right\}. \quad (3)$$

The n denotes the number of wave modes and k_n, c_n are the wavenumber and wave phase speed respectively. The subscript L denotes the variables at the bottom of stratosphere. N here is the buoyancy frequency and μ is thermal dissipation rate of waves. This GWF parameterization scheme (Eq. 3) can be derived adopting the local Boussinesq approximation and WKB approximation (Plumb [1975]). The integral in the exponential term further indicates that the forcing at height z depends on the vertical structure of the mean flow between z_L and z , implying intrinsic nonlocal dynamics of the background winds.

Although the 1D-QBO model neglects the Coriolis effect, it provides a fundamental framework for simulating QBO-like wind fields with downward propagating phases. Figure 2 shows the numerical solution of the 1D-QBO model with the boundary conditions,

$$\bar{u}(z = 17 \text{ km}) = \bar{u}(z = 35 \text{ km}) = 0 \text{ m s}^{-1}. \quad (4)$$

The parameters in equation (3) are set to $n = 2$ and $c_1 = -c_2 = 32 \text{ m} \cdot \text{s}^{-1}$, following Plumb [1977]. The model produces a zonal wind structure that resembles the observations (Figure 1), but with a more symmetric amplitude distribution. The gravity wave forcing in Figure 2(a, b) is also consistent with critical-layer theory, demonstrating a narrow region of strong momentum deposition just below the level where the mean zonal wind reaches its extremum.

After Plumb's work, various gravity wave parameterization (GWP) schemes have been developed to represent unresolved wave-mean flow interactions, such as Lindzen [1981] and Alexander and Dunkerton [1999]. These schemes have been adopted in the GCMs (Garfinkel et al. [2022]). Nevertheless,




Figure 2: The numerical solution of the 1D-QBO model. (a, b) the zonal wind $\bar{u}$ (blue line) and gravity wave forcing (GWF, green line) at day 200 and day 500 respectively; (c) the zonal wind profile over time.

the 1D-QBO model captures the essential wave-mean flow interactions responsible for the QBO and is easy to implement, so we employ Plumb’s model to generate training data for FNO models to learn. The basic parameters of the 1D-QBO model are shown in the Table 1.

Parameters	Quantity
buoyancy frequency N	$2.16 \times 10^{-2} s^{-1}$
diffusion rate κ	$0.3 m^2 s^{-1}$
n	2
c_1	$32 m s^{-1}$
c_2	$-32 m s^{-1}$
wavenumber k	$2\pi/40000 km^{-1}$
wave dissipation rate μ	$1.2 \times 10^{-6} s^{-1}$

Table 1: Parameters for 1D-QBO model.

2.2 Fourier Neural Operator

Fourier Neural Operator is a kind of Neural Operator which is designed to solve PDEs. Instead of learning maps $\mathbb{R}^n \rightarrow \mathbb{R}^m$ like most traditional neural networks, neural operators aim to learning maps in function spaces, i.e. learning the operator

$$\mathcal{G} : v(x) \rightarrow u(x), \quad (5)$$

where $v(x)$ is the input and $u(x)$ is the output (Kovachki et al. [2023]). In a classical PDE problem, $v(x)$ can be initial conditions, boundary conditions or forcing and $u(x)$ is the solution. Thus, the neural operators are often resolution-invariant because a function can be represented under different sampling schemes.

In order to guarantee the differentiability, The framework of neural operators is similar to neural networks,

$$u = (K_l \circ \sigma_l \circ \dots \circ \sigma_1 \circ K_1)(v), \quad (6)$$

where K_i are linear transforms and σ_i are activation functions. For FNO, K_i are Fourier layers, containing a Fourier transform and an inverse Fourier transform. The structure of a Fourier layer is demonstrated in Figure 3 (Li et al. [2020]). The input function $v(x)$ is first transformed into the frequency domain by a Fourier transform. A linear transformation R is then applied to the Fourier coefficients. The linear transforms in Fourier space are convolutions in real space, so FNOs resemble CNNs but with full domain extent, compared to the finite kernel size in CNNs. In practice, the spectrum is always truncated to avoid high frequency oscillation, which means the corresponding rows

(or columns) of matrix R are set to zero. After that, the result $R\mathcal{F}(x)$ is mapped back to the original space through an inverse Fourier transform. Finally, following the idea of residual connections, a linear transformation of the input function, $Wv(x)$, is added to the output of the Fourier layer. Since the derivatives of the Fourier transform and its inversion can both be expressed analytically, the Fourier layer is differentiable. The overall architecture of the Fourier layer can be written as,

$$k(v) = \mathcal{F}^{-1}(R\mathcal{F}(v)) + Wv, \quad (7)$$

where R and W are the learnable parameters.




Figure 3: The architecture of the Fourier layer (Li et al. [2020], Figure 2b).

2.3 Model Architecture




Figure 4: Model Architecture.

The architecture of the model is illustrated in Figure 4. The neural operator takes the zonal wind $\bar{u}^t$ at time t as input and predicts the corresponding GWF $\hat{f}^t$. First, the input is projected into a higher-dimensional feature space through a linear embedding layer. The tensor in feature space is then processed by a series of Fourier layers, each followed by an activation function. The output of the last Fourier layer is mapped to the target dimension through a fully connected layer.

The model is trained offline. The feature is $\bar{u}$ at every time step generated by the 1D-QBO model and the corresponding labels are given by $f^t = \text{GWF}(\bar{u}^t)$. These data are stochastically divided into the training set and the validation set. The order in each set is also randomized. The model performance is tested online. The predicted $\hat{f}^t$ is given to the numerical model to substitute $\text{GWF}(\bar{u}, z)$ in equation 3. If the model is denoted by $\mathcal{M}$, the mathematical expression of online emulation is

$$\frac{\partial \bar{u}}{\partial t} + w \frac{\partial \bar{u}}{\partial z} - \kappa \frac{\partial^2 \bar{u}}{\partial z^2} = \mathcal{M}(\bar{u}). \quad (8)$$

To compare the performance between different models, we also consider a fully connected neural network, i.e., multilayer perceptron (MLP). The deep MLP consists of eight hidden layers in order to extract enough features from inputs. The parameters defined in the two models are shown in table 2.

3 Result

3.1 Performance under Deterministic Training Data

We first consider a simple task in which the training datasets are generated without any noise. Figure 5 presents online testing results of the FNO and MLP models. The FNO model used here

	FNO	MLP
Number of layers	4	8
Learning rate	0.0001	0.1
Degrees of freedom	60012	2393978
Truncation mode	9	N/A
Channel	37	213 to 639

Table 2: Parameters for FNO and MLP.




Figure 5: Online emulations. (a) FNO; (b) MLP.

contains four Fourier layers. Both simulations are stable, consistent with the results in [Pahlavan et al. \[2025\]](#) because they take the full vertical profile as their receptive fields. However, the oscillation period generated by the MLP is shorter than that generated by the FNO: the MLP period is about 25.53 months, while FNO period is about 27.07 months, which is closer to the numerical model and real atmospheric phenomena.

Figure 6 demonstrates the offline and online GWF predictions provided by the two models. Although the two models perform excellently offline, the FNO model generates much smoother GWF, which contributes to its stable online performance. In contrast, the MLP predictions exhibit small fluctuations around $\text{GWF} \approx 0$. This instability exerts non-physical momentum flux onto the mean flow and thus causes the mean flow to be accelerated prematurely, which may shorten the oscillation period.

The improved period-stability of the FNO can possibly be understood from the structure of Fourier layers. As discussed in Section 2.2, the FNO retains only the first several low-wavenumber Fourier modes. This spectral truncation effectively suppresses high frequency fluctuations. In traditional neural networks (such as the MLP here), such fluctuations typically arise when the model encounters the inputs that are not well represented in the training dataset. Specifically, when the true forcing is close to $\text{GWF} \approx 0$, the training data contain many exact or near-zero values. Since the models approximate this value through nonlinear combinations of parameters and inputs, small prediction errors occur when new inputs are presented, leading to non-physical oscillations around zero. By contrast, although the Fourier layer appears to impose a rigid restriction by retaining only a limited set of spectral modes, this constraint prevents small prediction errors from being directly projected onto the forcing field. This mechanism is similar to a low pass filter. Consequently, the predicted GWF remains close to zero when the GWF is weak, avoiding artificial perturbations to the mean-flow acceleration and possibly contributing to a more stable oscillation period.

3.2 Performance under Noisy Training Data

In the real atmosphere (or oceans), the eddy forcing acting on the mean flow is inherently noisy due to unresolved processes and broadband wave spectra ([Berner et al. \[2017\]](#)). This noise will affect the stability and accuracy of both numerical models ([Hodyss et al. \[2013\]](#)) and ML models ([Wikner et al. \[2024\]](#)). Therefore, it is essential to assess the performance and generalization ability of data-driven models when trained on noisy datasets. From the ML perspective, this is also a robustness test for the application of FNO in GWP.




Figure 6: GWF prediction without noise. (a) Online performances; (b) Offline performances.

In this study, two different approaches are adopted to incorporate noise into the training data. In the first experiment, training data are generated from the deterministic numerical model, and additive white noise is subsequently imposed on the resulting dataset. This setup allows for an assessment of the model robustness when learning a deterministic physical process. Figure 7 shows the corresponding results. The online performance demonstrates that the FNO still produces a stable period, while the period generated by the MLP is much shorter than that shown in Figure 6.

Second, stochasticity is introduced directly into the gravity-wave forcing by perturbing the phase speeds of the waves, which is more similar to the real stratosphere (Holton and Lindzen [1972]). The numerical model is then driven by this stochastic forcing to generate training data. Instead of focusing solely on the impact of noise on the training process from a ML perspective, this experiment aims to assess whether the models can capture the underlying physical mechanisms when stochasticity is introduced to the governing dynamics. In Figure 8, both the FNO and MLP fail to reproduce the exact oscillation period. Nevertheless, the FNO yields a period that is substantially closer to the reference.

The potential failure of the FNO in the second task may stem from this tendency to learn the numerical average of stochastic GWF. However, the actual underlying dynamical evolution is not a simple linear superposition of these forced states. For example, the dynamical response to a fast-wave forcing followed by a slow-wave forcing, or the reverse, differs significantly from the response to the simple average of these two forcing terms. This implies that offline training is not sufficient to capture the nonlinear effect of stochasticity on long-term evolution, thereby underscoring the importance of online training to maintain physical consistency in nonlinear systems (Pahlavan et al. [2024]). In online training, the loss function contains both neural networks and the numerical model.

3.3 Performance of FNOs with Different Model Complexities

Although the above experiments demonstrate that the FNO has the potential to learn subgrid physics in geophysical dynamics, its computational cost remains a practical concern. When trained on the CPU of a personal laptop, the FNO requires approximately 90–120 minutes, whereas the MLP completes training within 15–30 minutes under comparable settings.

To assess the tradeoff between model complexity and performance, we examine FNO architectures with different numbers of Fourier layers. To eliminate the influence of stochastic noise for simplicity, we use deterministic forcing as a baseline. All models are trained for 200 epochs, which is sufficient to ensure convergence. The results indicate that models with insufficient complexity fail to produce stable online predictions. The FNO with one Fourier layer leads to rapid numerical divergence, while the one




Figure 7: GWF prediction with noisy training data. (a) Online performances; (b) Offline performances.




Figure 8: GWF prediction with stochastic model data. (a) Online performances; (b) Offline performances.

with two Fourier layers can generate a QBO-like structure with inconsistent periods at first, but also becomes unstable eventually. Once the number of Fourier layers reaches three or higher (results with four layers can be seen in Section 3.1), the model is able to generate stable oscillations with a period of approximately 28 months.

In this section, we show that while the Fourier layer can extract the physical features in the QBO data, the model’s ability is still restricted by the total number of parameters and depth. More parameters and deeper architectures provide the model more degrees of freedom to tackle the unseen input and enhance stability (Belkin et al. [2019]). In this specific task, a minimum of three Fourier layers with more than 40000 parameters is required to achieve a stable climate state in online emulations, although that is still much lower than millions of parameters in the MLP.

4 Conclusions and Discussions

In this study, we have reviewed the background of sub-grid scale turbulence parameterizations, moving from their necessity in geophysical models to the cutting-edge AI-based methodologies. The




Figure 9: Online performances of the models with different complexity. (a) one Fourier Layer; (b) two Fourier Layers; (c) three Fourier Layers.

rapid development of AI models has demonstrated remarkable efficacy in capturing complex fluid dynamic processes. Among these, the Fourier Neural Operator (FNO), a novel architecture, has achieved superior performance in global weather forecasting.

To evaluate the FNO’s capability in learning sub-grid physics, we employ a 1D Quasi-Biennial Oscillation (QBO) model as a testbed. The primary focus of this assessment is the model’s performance in online testing to see if it could faithfully reproduce the authentic QBO cycle. Previous studies and our findings all suggest that the offline accuracy of a model does not always translate to online stability. Although naive models can give relatively accurate results offline, their online performance is not necessarily unstable.

Under deterministic and noiseless conditions, the FNO demonstrates remarkable precision in maintaining the QBO period with noticeably fewer degrees of freedom compared to the MLP. While the MLP often suffers from significant phase drifts and subtle non-physical oscillations, the FNO successfully captures the sharp gradients associated with critical layer absorption. When additive white noise was introduced to the training data, the FNO exhibited robust generalization, whereas the MLP’s online performance degraded, resulting in a significantly shortened oscillation period. In experiments involving stochastic gravity-wave forcing, where the wave spectrum is broader, similar to that of the real stratosphere, both the FNO and the MLP fail to generate an accurate oscillation period, although the MLP’s period is much more inconsistent with the model truth. These results, on the one hand, illustrate that the FNO shows enhanced robustness for this task, learning Gravity Wave Forcing (GWF) in the 1D-QBO model. On the other hand, they imply that if we want to improve the generalization ability of the FNO, we need to make certain adjustments to the model structure, parameter size or training strategies. One potential strategy that has been widely used in eddy parameterization is online training, where the numerical model error is also taken into account in the training loss and thus the back propagation process.

Although the FNO can effectively learn the GWF in 1D-QBO model with a small set of parameters, the efficiency is another important indicator to evaluate the quality of a certain AI algorithm. Architectural sensitivity analysis revealed that the success of the FNO highly depends on its depth and complexity. We found that a minimum of three Fourier layers is essential to sustain a stable climate state in online emulations. With a parameter size of approximately 20,000, the FNO outperformed much larger MLP structures, further demonstrating its superiority to MLP architectures in gravity wave parameterization.

The outstanding performance of the FNO indicates that its architecture may in principle be able

to discover the underlying physics in QBO. We notice the small fluctuation in the MLP’s prediction and contrast it with the smooth profile generated by the FNO. We qualitatively attribute this stability to the spectral truncation inherent in the FNO architecture. This mechanism effectively suppresses non-physical noise introduced by the neural networks/operators, thereby enhancing the physical consistency of the long-term emulation.

Despite these conclusions, the physical consistency of FNO-based parameterizations still need to be investigated quantitatively. In this study, we focus on the accuracy of QBO periods as an indicator. However, a critical question remains whether the neural networks/operators can obey the physical conservation laws and whether the superior performance of the FNO can be explained from this aspect. The violations of these natural rules have been widely found in AI weather prediction (Bonavita [2024]), and the performance was significantly enhanced by improving this aspect (Sha et al. [2025]). Besides, some works also computed the probability density function (PDF) of GWF prediction to examine the physical consistency (Pahlavan et al. [2024]). To improve the physical consistency and preserve the advantages of neural operators, an alternative approach is Physics-Informed Neural Operator (PINO), which integrates data and physical constraints into the training session together (Li et al. [2024]). This strategy has been adopted in FNOs and shows better generalization ability in modeling chaotic Kolmogorov flows than traditional FNO and Physics-Informed Neural Networks (PINNs).

The internal architecture of the FNO offers a pathway for enhancing the explainability of the learned parameterization. As we mentioned in Section 2.2, the FNO shares structural similarities with CNNs. Previous studies utilizing CNNs for QBO parameterization have demonstrated that convolutional layers effectively function as multi-frequency filters and examined the Fourier spectra of kernels in each layer (Pahlavan et al. [2024]). The similar tests have the potential to be adopted on the FNO model here in order to check whether it preforms like a low-pass filter. Furthermore, the FNO transcends the limitations of traditional CNNs through its discretization-invariance, which is a unique property of neural operators (Li et al. [2020]). This facilitates further exploration of QBO phase sensitivity to grid resolution, a topic that has been studied in conventional numerical simulations (Giorgetta et al. [2002]).

Finally, the QBO remains a challenging signal to faithfully reproduce in General Circulation Models (GCMs) (Garfinkel et al. [2022]). However, its accurate representation is very important, given its profound impact on a wide range of atmospheric phenomena and global teleconnections (Anstey et al. [2022b]). Bridging the gap from this idealized 1D testbed to embedding FNO-based parameterizations into complex 3D GCMs remains the ultimate task and still has a long way to go. Such integrations will be critical to verify whether the spectral advantages identified in this study can effectively mitigate systematic biases in tropical stratospheric modeling and enhance the overall fidelity of global climate predictions.

References

- M. J. Alexander and T. J. Dunkerton. A spectral parameterization of mean-flow forcing due to breaking gravity waves. *Journal of the Atmospheric Sciences*, 56(24):4167–4182, 1999.
- J. A. Anstey, S. M. Osprey, J. Alexander, et al. Impacts, processes and projections of the quasi-biennial oscillation. *Nature Reviews Earth & Environment*, 3(9):588–603, 2022a.
- J. A. Anstey, I. R. Simpson, J. H. Richter, et al. Teleconnections of the quasi-biennial oscillation in a multi-model ensemble of qbo-resolving models. *Quarterly Journal of the Royal Meteorological Society*, 148(744):1568–1592, 2022b.
- M. P. Baldwin, L. J. Gray, T. J. Dunkerton, et al. The quasi-biennial oscillation. *Reviews of Geophysics*, 39(2):179–229, 2001.
- P. Bauer, A. Thorpe, and G. Brunet. The quiet revolution of numerical weather prediction. *Nature*, 525(7567):47–55, 2015.

- M. Belkin, D. Hsu, S. Ma, and S. Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854, 2019.
- J. Berner, U. Achatz, L. Batte, et al. Stochastic parameterization: Toward a new view of weather and climate models. *Bulletin of the American Meteorological Society*, 98(3):565–588, 2017.
- T. Beucler, M. Pritchard, S. Rasp, et al. Enforcing analytic constraints in neural networks emulating physical systems. *Physical review letters*, 126(9):098302, 2021.
- M. Bonavita. On some limitations of current machine learning weather prediction models. *Geophysical Research Letters*, 51(12):e2023GL107377, 2024.
- J. R. Booker and F. P. Bretherton. The critical layer for internal gravity waves in a shear flow. *Journal of fluid mechanics*, 27(3):513–539, 1967.
- C. O. de Burgh-Day and T. Leeuwenburg. Machine learning for numerical weather and climate modelling: a review. *Geoscientific Model Development*, 16(22):6433–6477, 2023.
- C. A. De Moura and C. S. Kubrusly. The courant–friedrichs–lewy (cfl) condition. *AMC*, 10(12):45–90, 2013.
- B. Fox-Kemper, A. Adcroft, C. W. Böning, et al. Challenges and prospects in ocean circulation models. *Frontiers in Marine Science*, 6:65, 2019.
- C. I. Garfinkel, E. P. Gerber, O. Shamir, et al. A qbo cookbook: Sensitivity of the quasi-biennial oscillation to resolution, resolved waves, and parameterized gravity waves. *Journal of Advances in Modeling Earth Systems*, 14(3):e2021MS002568, 2022.
- M. A. Giorgetta, E. Manzini, and E. Roeckner. Forcing of the quasi-biennial oscillation from a broad spectrum of atmospheric waves. *Geophysical Research Letters*, 29(8):86–1, 2002.
- R. Hallberg. Using a resolution function to regulate parameterizations of oceanic mesoscale eddy effects. *Ocean Modelling*, 72:92–103, 2013.
- H. T. Hewitt, M. Roberts, P. Mathiot, et al. Resolving and parameterising the ocean mesoscale in earth system models. *Current Climate Change Reports*, 6(4):137–152, 2020.
- D. Hodyss, K. C. Viner, A. Reinecke, et al. The impact of noisy physics on the stability and accuracy of physics–dynamics coupling. *Monthly weather review*, 141(12):4470–4486, 2013.
- J. R. Holton and R. S. Lindzen. An updated theory for the quasi-biennial cycle of the tropical stratosphere. *Journal of Atmospheric Sciences*, 29(6):1076–1080, 1972.
- N. Kovachki, Z. Li, B. Liu, et al. Neural operator: Learning maps between function spaces with applications to pdes. *Journal of Machine Learning Research*, 24(89):1–97, 2023.
- Z. Li, N. Kovachki, K. Azizzadenesheli, et al. Fourier neural operator for parametric partial differential equations. *arXiv preprint arXiv:2010.08895*, 2020.
- Z. Li, H. Zheng, N. Kovachki, et al. Physics-informed neural operator for learning partial differential equations. *ACM/IMS Journal of Data Science*, 1(3):1–27, 2024.
- R. S. Lindzen. Turbulence and stress owing to gravity wave and tidal breakdown. *Journal of Geophysical Research: Oceans*, 86(C10):9707–9714, 1981.
- R. S. Lindzen and J. R. Holton. A theory of the quasi-biennial oscillation. *Journal of Atmospheric Sciences*, 25(6):1095–1107, 1968.
- J. C. McWilliams. Statistical properties of decaying geostrophic turbulence. *Journal of Fluid Mechanics*, 198:199–230, 1989.
- J. D. Neelin. *Climate change and climate modeling*. Cambridge University Press, 2010.

- P. A. O’Gorman and J. G. Dwyer. Using machine learning to parameterize moist convection: Potential for modeling of climate, climate change, and extreme events. *Journal of Advances in Modeling Earth Systems*, 10(10):2548–2563, 2018.
- H. A. Pahlavan, P. Hassanzadeh, and M. J. Alexander. Explainable offline-online training of neural networks for parameterizations: A 1d gravity wave-qbo testbed in the small-data regime. *Geophysical Research Letters*, 51(2):e2023GL106324, 2024.
- H. A. Pahlavan, P. Hassanzadeh, and M. J. Alexander. On the importance of learning non-local dynamics for stable data-driven climate modeling: A 1d gravity wave-qbo testbed. *Geophysical Research Letters*, 52(10):e2024GL114136, 2025.
- R. A. Plumb. Momentum transport by the thermal tide in the stratosphere of venus. *Quarterly Journal of the Royal Meteorological Society*, 101(430):763–776, 1975.
- R. A. Plumb. The interaction of two internal waves with the mean flow: Implications for the theory of the quasi-biennial oscillation. *Journal of Atmospheric Sciences*, 34(12):1847–1858, 1977.
- S. Rasp, M. S. Pritchard, and P. Gentine. Deep learning to represent subgrid processes in climate models. *Proceedings of the national academy of sciences*, 115(39):9684–9689, 2018.
- J. H. Richter, J. A. Anstey, N. Butchart, et al. Progress in simulating the quasi-biennial oscillation in cmip models. *Journal of Geophysical Research: Atmospheres*, 125(8):e2019JD032362, 2020.
- Y. Sha, J. S. Schreck, W. Chapman, and D. J. Gagne. Improving ai weather prediction models using global mass and energy conservation schemes. *Journal of Advances in Modeling Earth Systems*, 17(11):e2025MS005138, 2025.
- M. Uliasz. *Subgrid-Scale Parameterizations*, pages 13–19. American Meteorological Society, Boston, MA, 1994. ISBN 978-1-935704-12-6.
- A. Wikner, J. Harvey, M. Girvan, et al. Stabilizing machine learning prediction of dynamics: Novel noise-inspired regularization tested with reservoir computing. *Neural Networks*, 170:94–110, 2024.
- F. E. Yan, H. Frezat, J. Le Sommer, J. Mak, and K. Otness. Adjoint-based online learning of two-layer quasi-geostrophic baroclinic turbulence. *Journal of Advances in Modeling Earth Systems*, 17(7):e2024MS004857, 2025.